

Importing the Libraries

```
In [1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns

from sklearn.model_selection import train_test_split
from sklearn.preprocessing import MinMaxScaler

from sklearn.linear_model import LogisticRegression
from sklearn.svm import SVC
from sklearn.ensemble import RandomForestClassifier
from sklearn.ensemble import GradientBoostingClassifier

from sklearn.metrics import confusion_matrix
from sklearn.metrics import classification_report
from sklearn.metrics import accuracy_score, recall_score, f1_score

from imblearn.over_sampling import RandomOverSampler
from imblearn.under_sampling import RandomUnderSampler
from collections import Counter

plt.style.use('fivethirtyeight')

import warnings
warnings.filterwarnings('ignore')
```

Reading & Exploring the dataset

```
In [2]: df = pd.read_csv('/kaggle/input/diabetes-dataset/diabetes.csv')
df
```

```
Out[2]:
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	DiabetesPedigreeFunction	Age	Outcome
0	6	148	72	35	0	33.6	0.627	50	
1	1	85	66	29	0	26.6	0.351	31	
2	8	183	64	0	0	23.3	0.672	32	
3	1	89	66	23	94	28.1	0.167	21	
4	0	137	40	35	168	43.1	2.288	33	
...
763	10	101	76	48	180	32.9	0.171	63	
764	2	122	70	27	0	36.8	0.340	27	
765	5	121	72	23	112	26.2	0.245	30	
766	1	126	60	0	0	30.1	0.349	47	
767	1	93	70	31	0	30.4	0.315	23	

768 rows × 9 columns

```
In [3]: df.isna().sum()
```

Pregnancies 0

```
Out[3]: Glucose      0
        BloodPressure 0
        SkinThickness 0
        Insulin       0
        BMI           0
        DiabetesPedigreeFunction 0
        Age          0
        Outcome       0
        dtype: int64
```

```
In [4]: df.duplicated().sum()
```

```
Out[4]: 0
```

```
In [5]: df.info()
```


```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 768 entries, 0 to 767
Data columns (total 9 columns):
#   Column              Non-Null Count  Dtype
---  -
0   Pregnancies         768 non-null   int64
1   Glucose             768 non-null   int64
2   BloodPressure       768 non-null   int64
3   SkinThickness       768 non-null   int64
4   Insulin             768 non-null   int64
5   BMI                 768 non-null   float64
6   DiabetesPedigreeFunction 768 non-null   float64
7   Age                 768 non-null   int64
8   Outcome             768 non-null   int64
dtypes: float64(2), int64(7)
memory usage: 54.1 KB
```

```
In [6]: # Showing the Correlations between columns
        df.corr()
```

```
Out[6]:
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	Diabetes
Pregnancies	1.000000	0.129459	0.141282	-0.081672	-0.073535	0.017683	
Glucose	0.129459	1.000000	0.152590	0.057328	0.331357	0.221071	
BloodPressure	0.141282	0.152590	1.000000	0.207371	0.088933	0.281805	
SkinThickness	-0.081672	0.057328	0.207371	1.000000	0.436783	0.392573	
Insulin	-0.073535	0.331357	0.088933	0.436783	1.000000	0.197859	
BMI	0.017683	0.221071	0.281805	0.392573	0.197859	1.000000	
DiabetesPedigreeFunction	-0.033523	0.137337	0.041265	0.183928	0.185071	0.140647	
Age	0.544341	0.263514	0.239528	-0.113970	-0.042163	0.036242	
Outcome	0.221898	0.466581	0.065068	0.074752	0.130548	0.292695	

```
In [7]: # Ploting the correlation
        plt.figure(figsize=(16, 10))
        sns.heatmap(df.corr(), annot=True, cmap='Blues')
        plt.title("Correlation Heatmap")
        plt.show()
```



We can notice that the Glucose column highly affect the outcome.
the BMI and Age columns may affect also the Outcome.
For the rest of the columns they rarely affect the Outcome.


EDA

```
In [8]: df['Outcome'].value_counts()
```

```
Out[8]: Outcome
0      500
1      268
Name: count, dtype: int64
```

```
In [9]: sns.countplot(x='Outcome', data = df)
```

```
Out[9]: <Axes: xlabel='Outcome', ylabel='count'>
```

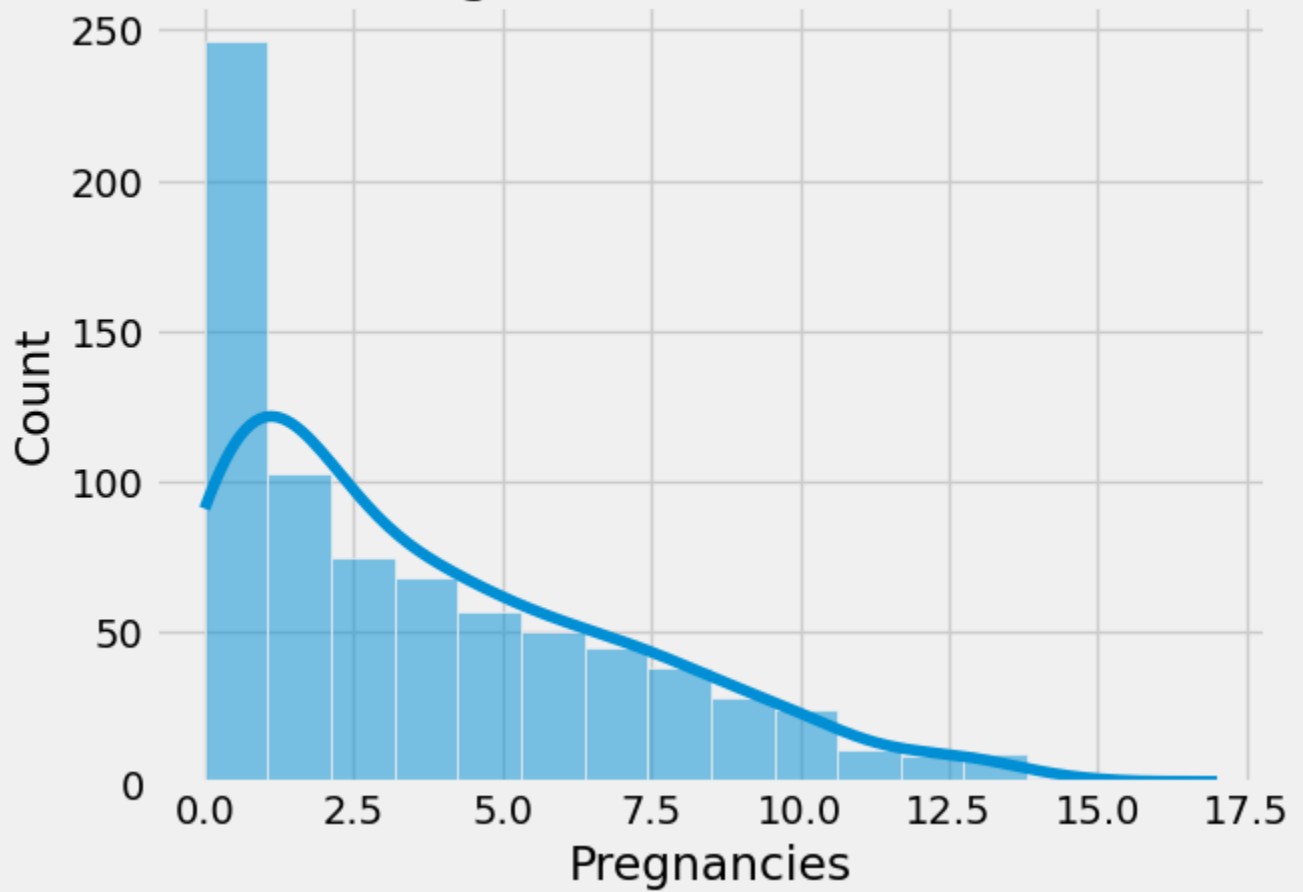


```
In [10]: numerical_columns = df.select_dtypes(exclude=object).columns.tolist()
```


```
In [11]: def num_cols_vis(col):  
    sns.histplot(data=df, x=col, kde=True)  
    plt.title(f'{col} Distribution')  
    plt.show()
```

```
In [12]: for col in numerical_columns:  
    num_cols_vis(col)
```

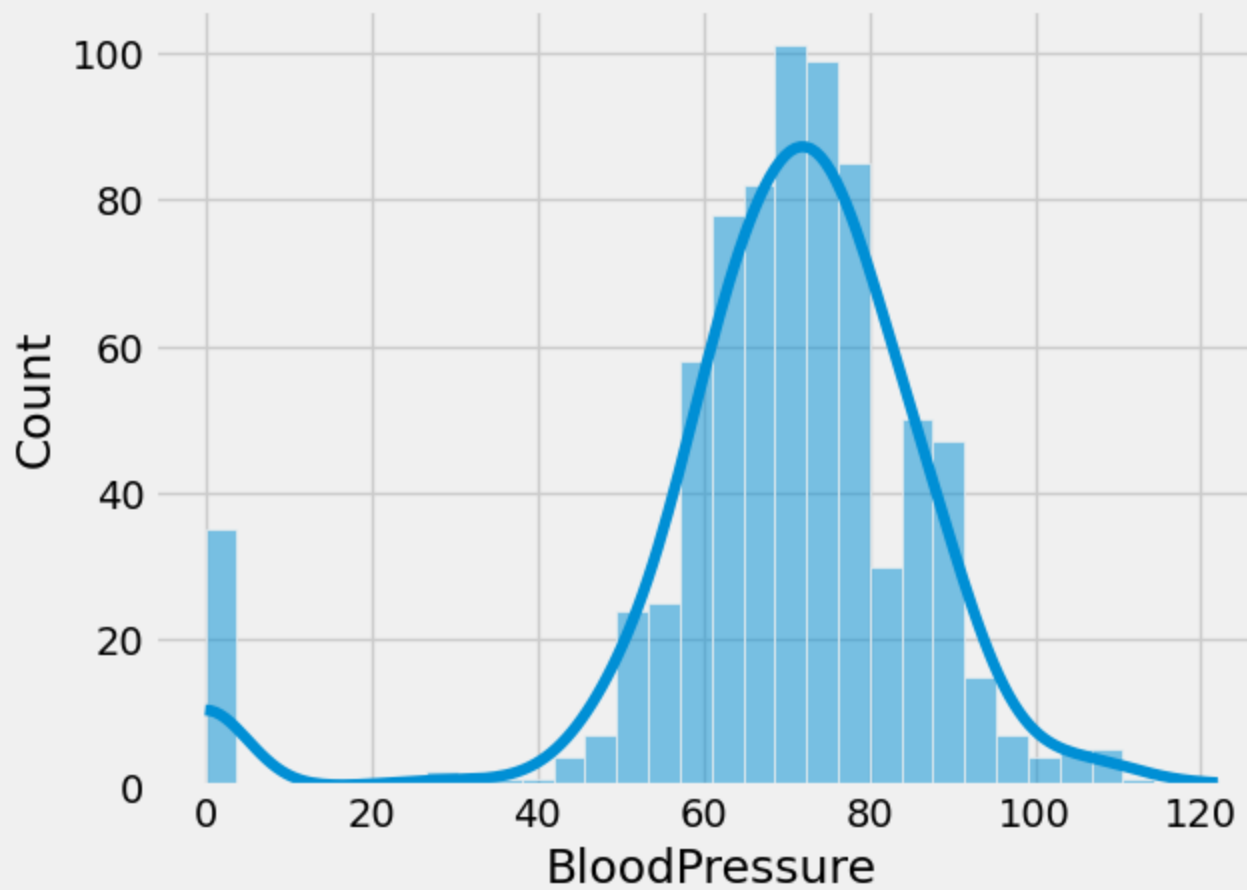
Pregnancies Distribution



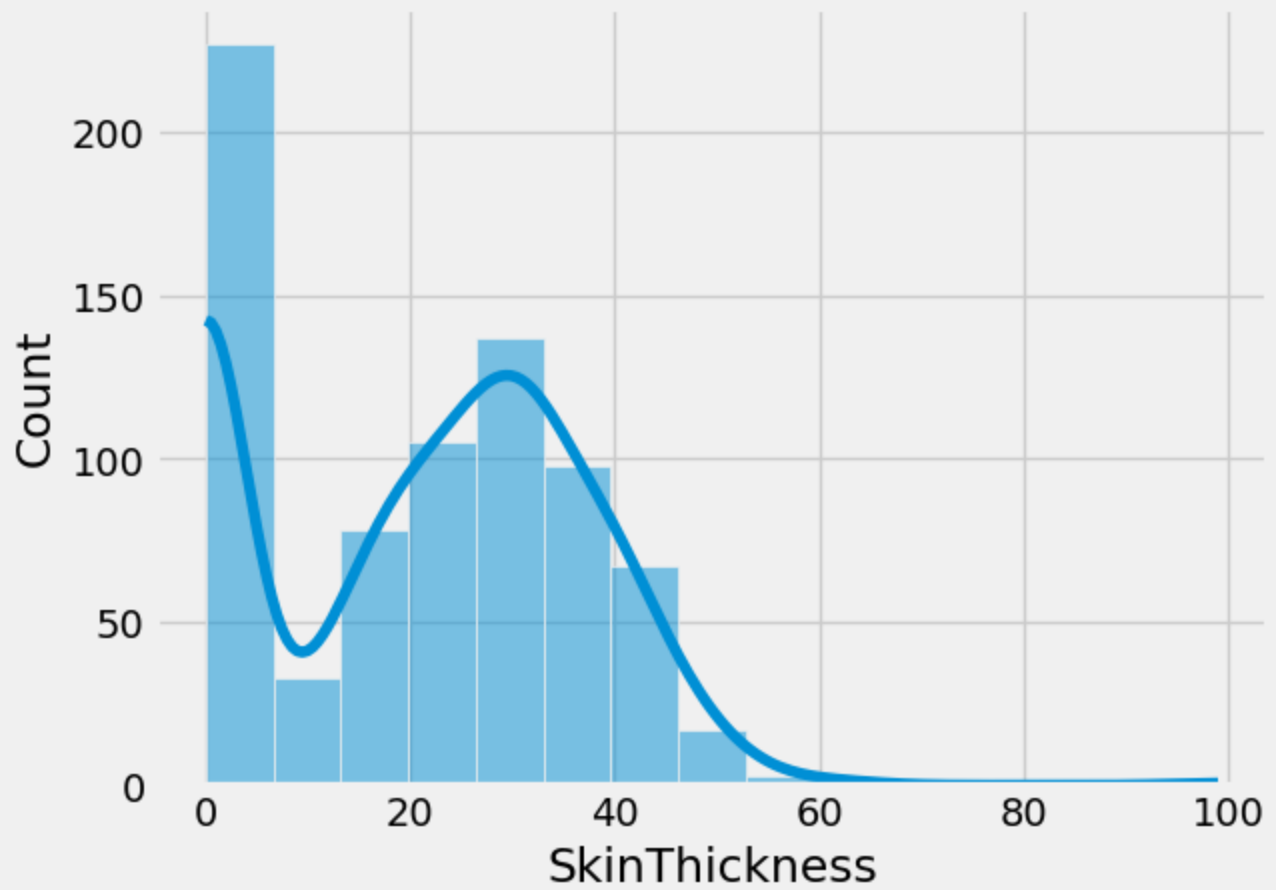
Glucose Distribution




BloodPressure Distribution



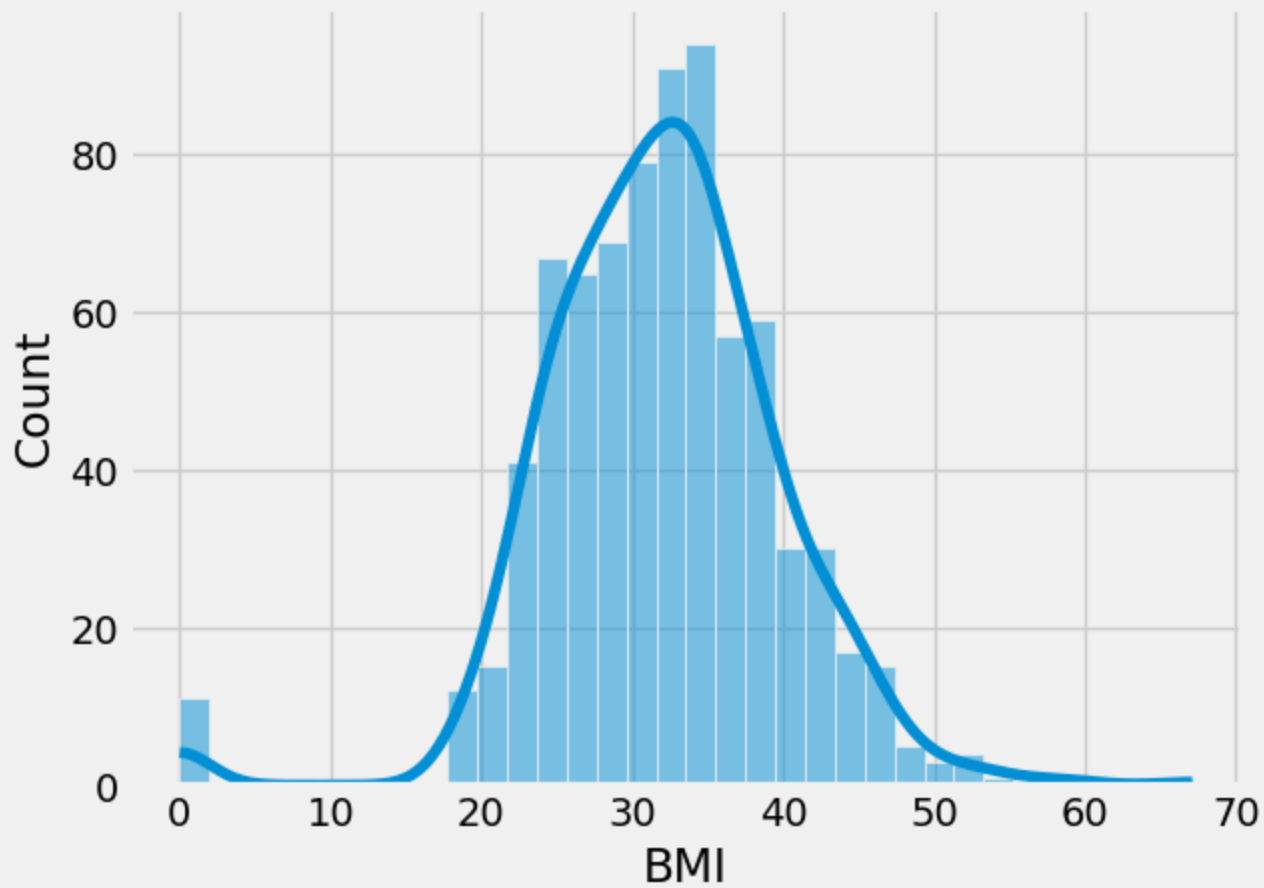
SkinThickness Distribution



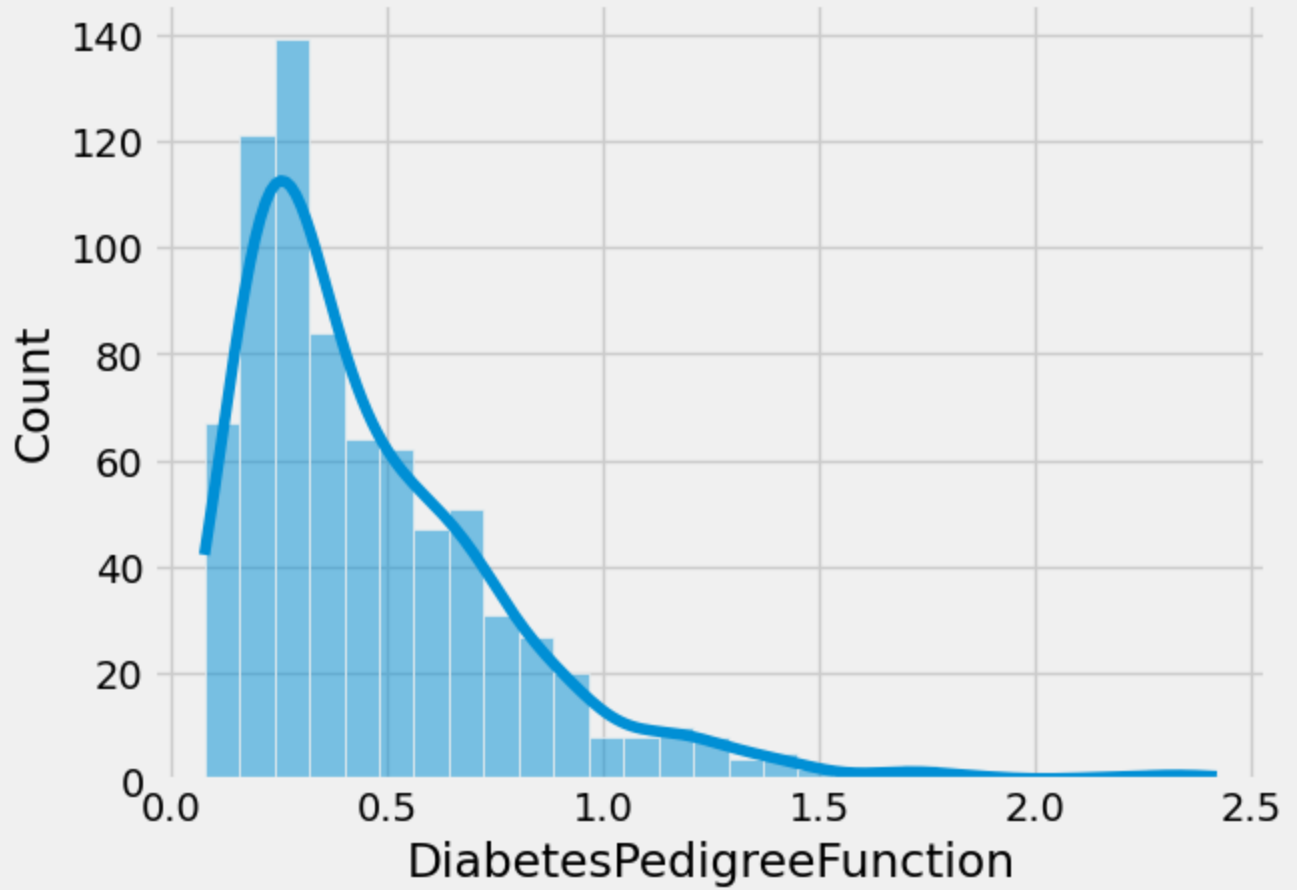
Insulin Distribution



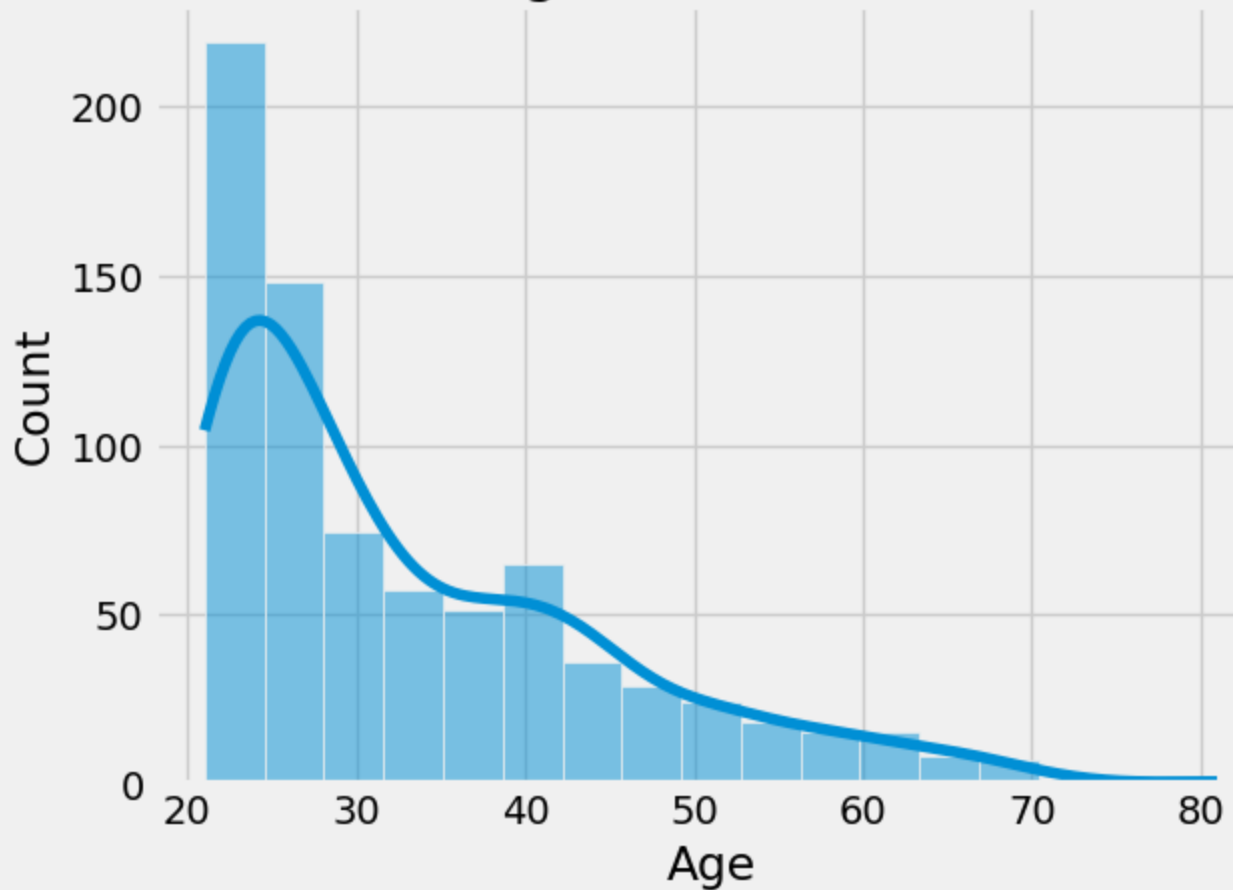
BMI Distribution

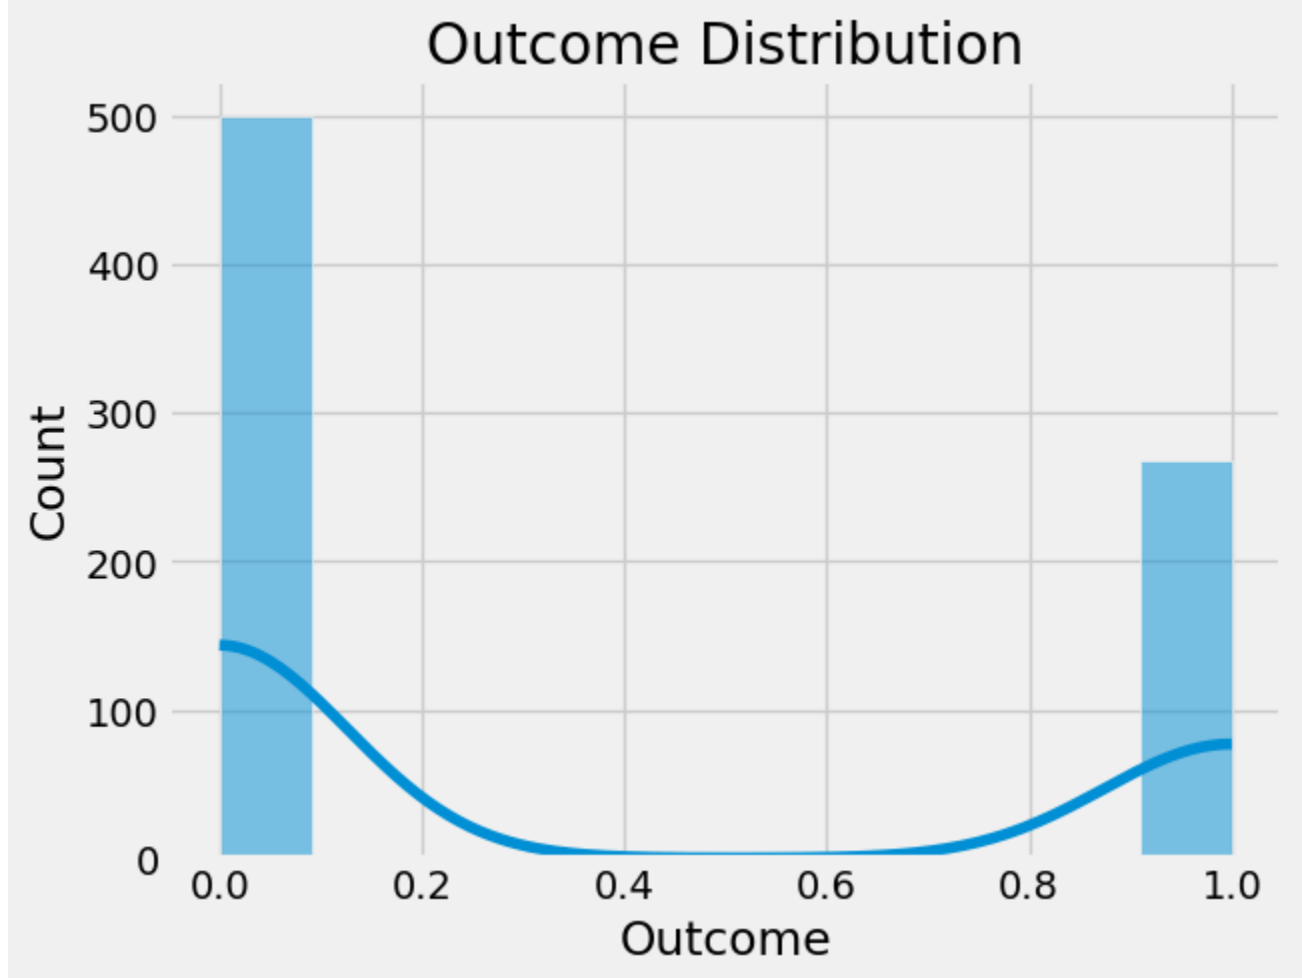


DiabetesPedigreeFunction Distribution



Age Distribution





Splitting the data

```
In [13]: x = df.drop('Outcome', axis=1)
         y = df['Outcome']
```

```
In [14]: rm = RandomOverSampler(random_state=41)
         x_res, y_res = rm.fit_resample(x, y)
```

```
In [15]: x_train, x_test, y_train, y_test = train_test_split(x_res, y_res, test_size=0.2)
```

Building the Models & Evaluation

```
In [16]: model_1 = LogisticRegression()
         model_2 = SVC()
         model_3 = RandomForestClassifier(n_estimators=100, class_weight='balanced')
         model_4 = GradientBoostingClassifier(n_estimators=1000)
```

```
In [17]: col = ['LogisticRegression', 'SVC', 'RandomForestClassifier', 'GradientBoostingClassifier']
         result_1 = []
         result_2 = []
         result_3 = []
```

```
In [18]: def cal(model):
         model.fit(x_train, y_train)
         pre = model.predict(x_test)
         accuracy = accuracy_score(pre, y_test)
         recall = recall_score(pre, y_test)
```

```

f1 = f1_score(pre,y_test)


result_1.append(accuracy)
result_2.append(recall)
result_3.append(f1)

sns.heatmap(confusion_matrix(pre,y_test),annot=True)
print(model)
print('Accuracy is: ',accuracy,'Recall is: ',recall,"F1 is: ",f1)
cal(model_1)

```

```
LogisticRegression()
```


```
Accuracy is: 0.745 Recall is: 0.75 F1 is: 0.7213114754098362
```



```
In [19]: cal(model_2)
```


```
SVC()
```

```
Accuracy is: 0.725 Recall is: 0.7380952380952381 F1 is: 0.6927374301675978
```




In [20]: `cal(model_3)`

```
RandomForestClassifier(class_weight='balanced')  
Accuracy is: 0.87 Recall is: 0.8349514563106796 F1 is: 0.8686868686868687
```



In [21]: `cal(model_4)`

```
GradientBoostingClassifier(n_estimators=1000)  
Accuracy is: 0.87 Recall is: 0.8415841584158416 F1 is: 0.8673469387755102
```



```
In [22]: final_result = pd.DataFrame({"Algorithm":col , 'Accuarcy':result_1,"recall":result_2,"F1_
final_result
```

Out[22]:

	Algorithm	Accuarcy	recall	F1_score
0	LogisticRegression	0.745	0.750000	0.721311
1	SVC	0.725	0.738095	0.692737
2	RandomForestClassifier	0.870	0.834951	0.868687
3	GradientBoostingClassifier	0.870	0.841584	0.867347

```
In [23]: # Performance Comparison of Classification Metrics Across Algorithms
fig,ax = plt.subplots(figsize=(15,5))
plt.plot(final_result.Algorithm,result_1,label="Accuracy")
plt.plot(final_result.Algorithm,result_2,label="Recall")
plt.plot(final_result.Algorithm,result_3,label="F1_Score")
plt.legend()
plt.show()
```

